


Co-occurrence counts on test set
(diagonal = correctly classified; off-diagonal = where the model got confused)



Per-class precision / recall / F1 (binary one-vs-rest)
Same F1 numbers as the main results table

